

Pig Script 1 :

```
##### The following is the Pig script for
Problem 1: Record lookup
#####
##### This script is to be stored as Pig1.pig
#####
#####
##### The script picks up input data from
directory - spark assignment created under HDFS on which root user has
Owner access #####
##### The script is run from Pig grunt , part of
Cloudera setup on AWS VM - CDH install
#####
##### Since, the script is being run from a directory -
Pig(created under root) , we could avoid giving absolute input path i.e.
user/root/.... ###
##### The output is being written to a directory
- pigoutput in root directory
#####
##### Document followed - How to run Pig on AWS
EC2 from Module 6 Session1
#####
```

```
data = LOAD 'spark_assignment/input_dataset/yellow_tripdata_*' USING
PigStorage(',') AS (f1:int, f2:chararray, f3:chararray, f4:int,
f5:double, f6:int, f7:bytearray,
f8:int, f9:int, f10:int, f11:double, f12:double, f13:double, f14:double,
f15:double, f16:double, f17:double);
```

```
filtered = FILTER data BY f1 == 2 AND f2 == '2017-10-01 00:15:30' AND f3
== '2017-10-01 00:25:11' AND f4 == 1 AND f5 == 2.17 ;
```

```
STORE filtered INTO 'pigoutput/PigOutput1.out';
```

```
##### End of Pig1.pig
#####
```

Pig Script 2 :

```
##### The following is the Pig script for
Problem 2: Filter Records
#####
##### This script is to be stored as Pig2.pig
#####
#####
##### The script picks up input data from
directory - spark assignment created under HDFS on which root user has
Owner access #####
##### The script is run from Pig grunt , part of
Cloudera setup on AWS VM - CDH install
#####
##### Since, the script is being run from a directory -
Pig(created under root) , we could avoid giving absolute input path i.e.
user/root/.... ###
```

```
##### The output is being written to a directory
- pigoutput in root directory
#####
##### Document followed - How to run Pig on AWS
EC2 from Module 6 Session1
#####
```

```
data = LOAD 'spark_assignment/input_dataset/yellow_tripdata_*' USING
PigStorage(',') AS (f1:int, f2:chararray, f3:chararray, f4:int,
f5:double, f6:int, f7:bytearray,
f8:int, f9:int, f10:int, f11:double, f12:double, f13:double, f14:double,
f15:double, f16:double, f17:double);
```

```
filtered = FILTER data BY f6 == 4;
```

```
STORE filtered INTO 'pigoutput/PigOutput2.out';
```

```
##### End of Pig2.pig
#####
```

Pig Script 3 :

```
##### The following is the Pig script for
Problem 3: Group By , Count and Sort
#####
##### This script is to be stored as Pig3.pig
#####
#####
##### The script picks up input data from
directory - spark assignment created under HDFS on which root user has
Owner access #####
##### The script is run from Pig grunt , part of
Cloudera setup on AWS VM - CDH install
#####
##### Since, the script is being run from a directory -
Pig(created under root) , we could avoid giving absolute input path i.e.
user/root/.... ###
##### The output is being written to a directory
- pigoutput in root directory
#####
##### Document followed - How to run Pig on AWS
EC2 from Module 6 Session1
#####
```

```
data = LOAD 'spark_assignment/input_dataset/yellow_tripdata_*' USING
PigStorage(',') AS (VendorID,
tpep_pickup_datetime,
tpep_dropoff_datetime,
passenger_count,
trip_distance,
RatecodeID,
store_and_fwd_flag,
PULocationID,
DOLocationID,
payment_type,
fare_amount,
extra,
```

```

mta_tax,
tip_amount,
tolls_amount,
improvement_surcharge,
total_amount
);

data_not_null = FILTER data BY payment_type != '';
grouped_data = GROUP data_not_null by payment_type;
group_count = FOREACH grouped_data GENERATE group AS PAYMENT_TYPE,
COUNT_STAR(data_not_null) AS cnt;
sorted_count = ORDER group_count BY cnt;
result_without_header = FILTER sorted_count BY $0 != 'payment_type';
STORE result_without_header INTO 'pigoutput/PigOutput3.out';

##Line no 76 filters out all the non-null records
##Line no 77 groups by payment_type on not null dataset
##Line no 78 applies COUNT_STAR operation on non-null dataset so that
count of records corresponding to each payment_type could be found
##Line no 79 applies ORDER By on count obtained, in ascending order
##Line no 80 gets the records by removing the header row

##### End of Pig3.pig
#####

```

Spark Java Code for problem 1

```

##### The following is the Java code for Problem
1: Record lookup
#####
##### This code is to be stored as
AssignmentProblem1.java
#####
##
##### The program picks up input data from
directory - spark assignment created under HDFS on which root user has
Owner access #####
##### The script is run from Shell corresponding
to AWS VM - CDH install
#####
##### The script is being run from root, the
executable jar is placed in /home/ec2-user directory
#####
##### Hence, for input path , the absolute path
for input dataset is specified
#####
##### The output is being written to a directory
created under HDFS on which root user has Owner access
#####
##### References - Guidelines for Running Spark
Codes in EC2 Unix Box in project2
#####

```

```

package spark.assignment1;

//Importing all the required Spark packages

import org.apache.spark.SparkConf;

```

```

import org.apache.spark.api.java.JavaRDD;
import org.apache.spark.api.java.JavaSparkContext;
import org.apache.spark.api.java.function.Function;

public class AssignmentProblem1 {

    public static void main(String[] args) {

        //Instantiating the spark conf instance

        SparkConf conf = new
SparkConf().setAppName("MyFirstSparkProgram").setMaster("local[*]");
        System.out.println("Conf has been set");

        //Declaring a JavaRDD and initializing it to null

        JavaRDD<String> stringRDD = null;

        //Need a try catch block here so that any problem in
creating Spark Context
        // Or in reading the input file be handled gracefully

        try {

            JavaSparkContext sc = new JavaSparkContext(conf);
            System.out.println("Context has been
initialized");

            //Reading the input data set as a text file in
JavaRDD created
            stringRDD =
sc.textFile("/user/root/spark_assignment/input_dataset/yellow_tripdata*")
;
            System.out.println("File has been read
successfully");
        }catch(Exception e) {

            //returning the stack trace of Exception
encountered
            e.printStackTrace();
        }

        //Implementing filter transformation on JavaRDD which
has input data set

        JavaRDD<String> lookupRDD = stringRDD.filter(new
Function<String, Boolean>(){
            public Boolean call(String row) throws Exception {

                //Reading each record and splitting it by:
", "

                //Adding the split record to an array

                String[] rowValues = row.split(",");

                //Checking if the record under processing is
not null

                //and has length at least = 17
                if(rowValues!=null && rowValues.length>=17) {

```

```

//Below is the filter condition
as per Problem no 1
        if("2".equals(rowValues[0]) && "2017-
10-01 00:15:30".equals(rowValues[1]) && "2017-10-01
00:25:11".equals(rowValues[2]) && "1".equals(rowValues[3]) &&
"2.17".equals(rowValues[4])) {
                return true;
        } //Only records are returned in
lookupRDD which satisfy the filter condition
        }
        return false;
    }
}; //lambda function ends here

System.out.println("Writing the output now");

//Writing the output to a text file which will be saved
at location as given in command line - args[1]
lookupRDD.saveAsTextFile(args[1]);

System.out.println("File has been written with output");
}

}

```

```

##### End of AssignmentProblem1.java
#####
Spark Java Code for problem 2

```

```

##### The following is the Java code for Problem
2: Filter Records
#####
##### This code is to be stored as
AssignmentProblem2.java
#####
##
##### The program picks up input data from
directory - spark assignment created under HDFS on which root user has
Owner access #####
##### The script is run from Shell corresponding
to AWS VM - CDH install
#####
##### The script is being run from root, the
executable jar is placed in /home/ec2-user directory
#####
##### Hence, for input path , the absolute path
for input dataset is specified
#####
##### The output is being written to a directory
created under HDFS on which root user has Owner access
#####
##### References - Guidelines for Running Spark
Codes in EC2 Unix Box in project2
#####

```

```

package spark.assignment2;

//Importing all the required Spark packages

import org.apache.spark.SparkConf;
import org.apache.spark.api.java.JavaRDD;
import org.apache.spark.api.java.JavaSparkContext;
import org.apache.spark.api.java.function.Function;

public class AssignmentProblem2 {

    public static void main(String[] args) {

        //Instantiating the spark conf instance

        SparkConf conf = new
SparkConf().setAppName("MySecondSparkProgram").setMaster("local[*]");
        System.out.println("Conf has been set");

        //Declaring a JavaRDD and initializing it to null

        JavaRDD<String> stringRDD = null;

        //Need a try catch block here so that any problem in creating
Spark Context
        // Or in reading the input file be handled gracefully

        try {

            JavaSparkContext sc = new JavaSparkContext(conf);
            System.out.println("Context has been initialized");

            //Reading the input data set as a text file in JavaRDD
created

            stringRDD =
sc.textFile("/user/root/spark_assignment/input_dataset/yellow_tripdata*")
;
            System.out.println("File has been read successfully");
        }catch(Exception e) {

            //returning the stack trace of Exception encountered
            e.printStackTrace();
        }

        //Implementing filter transformation on JavaRDD which has
input data set

        JavaRDD<String> filterRDD = stringRDD.filter(new
Function<String,Boolean>(){

            @Override
            public Boolean call(String row) throws Exception {

                //Reading each record and splitting it by: ","
                //Adding the split record to an array
                String[] rowValues = row.split(",");

                //Checking if the record under processing is not
null

```

```

        //and has length at least = 17

        if(rowValues!=null && rowValues.length>=17) {

            //Below is the filter condition as per
Problem no 2
            if("4".equals(rowValues[5])) {
                return true;
            }
        }
        return false;
    }
}

); //lambda function ends here
System.out.println("Writing the output now");

//Writing the output to a text file which will be saved at
location as given in command line - args[1]
filterRDD.saveAsTextFile(args[1]);

System.out.println("File has been written with output");

}

}

```

```

##### End of AssignmentProblem2.java
#####

```

Spark Java Code for problem 3

```

##### The following is the Java code for Problem
3: Group By , Count and Sort
#####
##### This code is to be stored as
AssignmentProblem2.java
#####
##
##### The program picks up input data from
directory - spark assignment created under HDFS on which root user has
Owner access #####
##### The script is run from Shell corresponding
to AWS VM - CDH install
#####
##### The script is being run from root, the
executable jar is placed in /home/ec2-user directory
#####
##### Hence, for input path , the absolute path
for input dataset is specified
#####
##### The output is being written to a directory
created under HDFS on which root user has Owner access
#####
##### References - Guidelines for Running Spark
Codes in EC2 Unix Box in project2
#####

```

```

package spark.assignment3;

//Importing all the required Spark packages

import org.apache.spark.SparkConf;
import org.apache.spark.api.java.JavaPairRDD;
import org.apache.spark.api.java.JavaRDD;
import org.apache.spark.api.java.JavaSparkContext;
import org.apache.spark.api.java.function.Function;

import scala.Tuple2;

public class AssignmentProblem3 {

    public static void main(String[] args) {

        //Instantiating the spark conf instance

        SparkConf conf = new
SparkConf().setAppName("MyThirdSparkProgram").setMaster("local[*]");
        System.out.println("Spark Conf has been set successfully");

        //Declaring a JavaRDD and initializing it to null

        JavaRDD<String> stringRDD = null;

        //Need a try catch block here so that any problem in creating
Spark Context
        // Or in reading the input file be handled gracefully

        try {

            JavaSparkContext sc = new JavaSparkContext(conf);
            System.out.println("Context has been initialized");

            //Reading the input data set as a text file in JavaRDD
created
            stringRDD =
sc.textFile("/user/root/spark_assignment/input_dataset/yellow_tripdata*")
;
            System.out.println("File has been read successfully");

        }catch(Exception e) {

            //returning the stack trace of Exception encountered
            e.printStackTrace();
        }

        //Implementing filter transformation on JavaRDD which has
input data set

        JavaRDD<String> resultantRDD = stringRDD.filter(new
Function<String,Boolean>(){
            @Override
            public Boolean call(String row) throws Exception {

                //Reading each record and splitting it by: ","
                //Adding the split record to an array

```



```

        String[] rowValues = row.split(",");

        //Checking if the record under processing is not
null
        //and has length at least = 17

        if(rowValues!=null && rowValues.length>=17) {

            //Filtering out the non-null rows and then
filtering out the header row
            if( (!rowValues[9].isEmpty())) {

                if(!("payment_type".equalsIgnoreCase(rowValues[9])))
                    return true;

                }

            }
            return false;

        }

    }; //lambda function ends here

    //Using map to Pair transformation to get key value pairs
    JavaPairRDD<String , Integer> pairRDD =
resultantRDD.mapToPair(

        //This is going to split the record and get value
of payment type column : key-> Payment_type , value -> 1
        x -> new Tuple2<String , Integer>(x.split(",")[9]
, 1)

    );

    //The reduce by key transformation is going to give the sum of
all records of a payment type
    JavaPairRDD<String , Integer> groupRDD = pairRDD.reduceByKey(
        (x,y) -> x+y
    );

    //The result will be sorted by key by default
    //We first swap the key and value so that result can be sorted
by value
    //We again swap back the result so that we have payment type
as key and sum of records of that payment type as value

    groupRDD = groupRDD.mapToPair(x ->
x.swap()).sortByKey().mapToPair(x -> x.swap());

    System.out.println("Writing the output now");

    //Writing the output to a text file which will be saved at
location as given in command line - args[1]
    groupRDD.saveAsTextFile(args[1]);
    System.out.println("File has been written with output");

    //groupRDD.foreach(x -> System.out.println(x._1 + ":" +
x._2));

    }

}

```

```
##### End of AssignmentProblem3.java
#####
```